Employee Attrition Prediction with Explainable AI, Transformer Models, and LDA Topic Modeling: An Interpretable Framework for Workforce Analytics

Kalla Praveen Balaji, Kannekanti Pavan Teja, N. Umasankari, and D. Geethanjali

Department of Computer Science and Engineering,
Sathyabama Institute of Science and Technology, Chennai, India

Abstract. Abstract. Employee attrition remains a serious organizational issue because it affects continuity, productivity, and workforce stability. Although machine learning models have improved the prediction of employee departure, their use in human resource decision-making is often constrained by limited interpretability. This study presents an integrated framework that combines Transformer-based learning for structured employee data with explainable artificial intelligence and topic modeling of employee feedback. Multi-head self-attention is used to capture complex interactions among workforce variables, while SHAP provides interpretable feature-level explanations at both individual and overall levels. In parallel, Latent Dirichlet Allocation is used to identify recurring themes in unstructured employee narratives, thereby enriching the interpretation of attrition risk. The framework is evaluated on benchmark employee attrition data and compared with common machine learning and neural baselines. Performance is assessed using accuracy, precision, recall, F1 score, and ROC-AUC, while equal attention is given to transparency and managerial usefulness. The study contributes an integrated framework for workforce analytics that balances predictive strength with interpretability.

Keywords: Employee attrition prediction, Transformer neural networks, Explainable artificial intelligence, SHAP interpretability, Topic modeling, Human resource analytics

1 Introduction

Employee attrition remains a major challenge for organizations because it increases recruitment and training costs, interrupts workflow continuity, and can weaken morale and knowledge retention. In knowledge-intensive environments, employee departure may also erode institutional memory and reduce operational resilience. As a result, retention has become a strategic issue rather than merely an administrative HR concern.

Conventional approaches to understanding attrition, such as exit interviews, surveys, and managerial judgement, are useful but largely reactive. They explain departure after it happens rather than supporting early intervention. With the growing availability of digitized HR records, predictive analytics has emerged as a practical way to identify employees who may be at risk before separation occurs.

Machine learning has contributed substantially to this transition. Models such as logistic regression, decision trees, support vector machines, random forests, and boosting-based methods have shown that structured HR variables can support effective attrition forecasting. However, these methods expose a central trade-off. Simpler models are easier to interpret but may not capture complex employee behaviour, whereas stronger predictive models often sacrifice transparency. In HR contexts, where outcomes affect people directly, explanation is as important as prediction.

Recent progress in deep learning offers new possibilities. Transformer-based models, originally developed for sequence tasks, have increasingly been adapted for tabular data. Their attention mechanisms help capture interactions among compensation, workload, satisfaction, career growth, managerial support, and tenure without extensive manual feature engineering. At the same time, explainable AI methods such as SHAP make it possible to interpret model outputs through feature-level contributions.

Another limitation in many attrition studies is the underuse of narrative employee feedback. Reviews, comments, and open-ended survey responses often reveal issues related to culture, recognition, stress, managerial climate, and development barriers that are not fully reflected in structured records. Topic modeling can expose these hidden patterns and provide richer context.

Motivated by these gaps, this study proposes an integrated framework that combines a Transformer-based tabular model, SHAP-based explanation, and topic modeling for employee feedback. The framework aims to improve predictive performance while retaining managerial interpretability. The remainder of the paper reviews relevant literature, presents the framework, describes the evaluation design, reports findings, and discusses implications.

2 Literature Review

This section reviews four connected streams relevant to the proposed study: machine learning for attrition prediction, explainable AI in HR analytics, Transformer-based tabular learning, and topic modeling for employee feedback.

2.1 Recent Employee Attrition Prediction Research

Employee attrition has been widely studied through supervised learning applied to structured HR data. Mansor *et al.* [7] showed that support vector machines, decision trees, and neural networks remain useful when preprocessing and class imbalance are handled appropriately. Raza *et al.* [11] compared multiple methods and highlighted variables such as income, age, and job level as important predictors. Chung *et al.* [2] further demonstrated that stacking-based ensembles can outperform individual classifiers. More recent work by Varkiani *et al.* [15] and Konar *et al.* [6] shifted attention from raw predictive performance toward explanatory determinants and explainability-aware modeling. Nandal *et al.* [8] also pointed out the baseline relevance of evaluating feature distributions carefully.

These studies confirm that data-driven forecasting is valuable, but most work still emphasizes conventional machine learning and ensemble methods. Comparatively fewer studies have explored deep architectures specifically designed to model complex interactions in structured employee data.

2.2 Explainable Artificial Intelligence in HR Analytics

Interpretability is especially important in HR analytics because predictions may influence decisions with direct human consequences. A model may achieve high accuracy yet still be unsuitable if managers cannot understand why a particular employee is flagged as high risk. Explainable AI addresses this issue by revealing how features contribute to predictions.

Das *et al.* [3] showed that feature-level explanation can strengthen attrition analysis by clarifying why particular employees are assessed as vulnerable. Sekaran and Sundaramurthy [12] similarly argued that useful attrition systems must identify major drivers rather than merely output class labels. Baydili and Tasci [1] positioned explainability as a design requirement in managerial decision support, while Konar *et al.* [6] showed that explainability can be combined with high-performing predictive frameworks. Together, these studies indicate that interpretability is no longer optional in workforce analytics.

2.3 Alternative Tabular Approaches

Transformer architectures have gained importance because attention mechanisms can model complex dependencies among inputs. Although they were first developed for language tasks, recent work has adapted them effectively for structured tabular problems. Gorishniy *et al.* [5] showed that carefully designed deep models, including Transformer variants, can compete with strong gradient-boosting baselines. Somepalli *et al.* [13] proposed SAINT, which uses attention across rows and columns and demonstrated the value of attention-based representation learning for tabular datasets.

This development is particularly relevant for attrition, where departure often reflects interacting conditions rather than isolated variables. Patterns such as overtime combined with dissatisfaction or long tenure with limited promotion are more naturally modeled through attention-based architectures than through simpler independent-feature assumptions.

2.4 Text Analytics and Topic Modeling for Employee Insights

Structured HR records do not always capture the full employee experience. Narrative data from reviews, comments, and open-ended responses often contains signals related to workplace climate, recognition, compensation fairness, and managerial quality. Putra and Puspitasari [9] showed that text analytics and topic modeling can surface actionable organizational issues from employee reviews. Ding *et al.* [4] extended this direction by combining structural topic modeling with supervised validation to derive more systematic employee insights.

These contributions suggest that textual analysis can complement structured attrition models by revealing concerns that remain hidden in quantitative variables alone.

2.5 HR Analytics and Strategic Retention Context

Broader reviews also emphasize the strategic role of analytics and AI in retention. Ravesangar and Narayanan [10] discussed the movement toward evidence-driven retention management, while Ubeda-Garcia *et al.* [14] examined wider tensions involving governance, knowledge, and organizational readiness in AI-enabled HR systems. These perspectives reinforce the need for attrition analytics that are predictive, interpretable, and responsibly deployable.

2.6 Research Gap

Three main gaps emerge. First, Transformer-style tabular learning remains less explored than classical machine learning in attrition studies. Second, explainability is often added after modeling rather than embedded in the framework itself. Third, structured HR records and unstructured employee feedback are commonly analyzed separately. The proposed study addresses these limitations by integrating Transformer-based tabular learning, SHAP-based explanation, and topic modeling in one unified framework.

3 Methodology

This section presents the proposed framework, including its architecture, data preparation pipeline, prediction model, explanation layer, topic-modeling component, and risk categorization logic.

3.1 Framework Architecture Overview

The framework consists of three connected layers: data processing, intelligence, and analytics visualization. The data processing layer prepares structured and unstructured inputs. The intelligence layer contains the Transformer-based predictor, the SHAP explanation engine, and the topic-modeling module. The analytics layer converts these outputs into interpretable risk profiles and decision-support summaries.

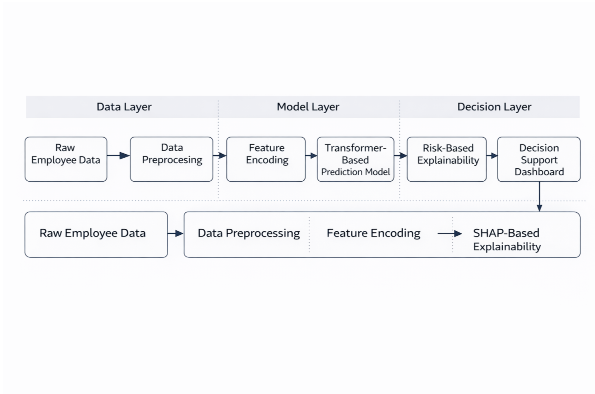

Fig. 1: System architecture of the proposed employee attrition prediction framework.

3.2 Data Processing Layer

The structured component may include demographic, compensation, job, satisfaction, performance, and tenure-related attributes. Numerical variables are standardized, categorical variables are encoded through learnable representations, and missing values are handled through appropriate imputation strategies. The unstructured component consists of employee feedback text, which is processed through tokenization, lowercasing, stop-word removal, and lemmatization before conversion to a document-term representation for topic modeling.

3.3 Transformer-Based Attrition Prediction Model

Let the employee feature vector be represented as

$$\mathbf{x} = [x_1, x_2, \ldots, x_n]^T, \tag{1}$$

where each component corresponds to a structured HR attribute such as age, income, job satisfaction, performance rating, or years at the company.

The input feature vector is transformed into an embedding representation and projected into query, key, and value matrices:

$$\mathbf{Q} = \mathbf{X}W_Q, \qquad \mathbf{K} = \mathbf{X}W_K, \qquad \mathbf{V} = \mathbf{X}W_V, \tag{2}$$

where W_Q, W_K, and W_V are learnable weight matrices.

The relationships among employee attributes are modeled using scaled dot-product self-attention:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right)\mathbf{V}, \tag{3}$$

where d_k denotes the dimensionality of the key vectors.

Multiple attention heads are then combined to capture diverse feature interactions:

$$\text{MultiHead}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{Concat}(\text{head}_1, \ldots, \text{head}_h)W_O, \tag{4}$$

where

$$\text{head}_i = \text{Attention}(\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i). \tag{5}$$

The resulting encoder representation is aggregated and passed through a sigmoid activation function to estimate the employee attrition probability:

$$\hat{y} = \sigma(\mathbf{W}_o \mathbf{z} + b), \tag{6}$$

where $\hat{y} \in [0, 1]$ represents the predicted probability of employee attrition.

3.4 SHAP Explainability Layer

SHAP is used to decompose predictions into feature-level contributions. At the local level, it shows which features push an individual employee toward or away from predicted attrition risk. At the global level, aggregated absolute SHAP values reveal the dominant drivers of attrition across the population.

3.5 Topic Modeling for Employee Feedback

To complement structured prediction, the framework includes a Latent Dirichlet Allocation module. Each document is modeled as a mixture of latent topics, with each topic defined as a distribution over words. This enables discovery of recurring concerns such as compensation dissatisfaction, limited growth, management quality, and work-life balance. Topic prevalence may be used as auxiliary input or as contextual interpretation alongside structured predictions.

3.6 Risk Classification Engine

For practical decision support, predicted attrition probability is converted into four risk categories:

- **High Risk** ($\hat{y} \geq 0.70$): immediate action
- **Elevated Risk** ($0.40 \leq \hat{y} < 0.70$): proactive intervention
- **Moderate Risk** ($0.20 \leq \hat{y} < 0.40$): monitoring
- **Low Risk** ($\hat{y} < 0.20$): routine management

Each risk profile can be accompanied by leading SHAP-based drivers and relevant topic cues.

4 Experimental Setup

This section outlines the dataset, evaluation strategy, baseline models, and implementation details used to assess the framework. The focus is on ensuring a transparent and reproducible evaluation design aligned with prior research in employee attrition analytics.

Algorithm 1 Integrated attrition analytics workflow

Require: Structured employee data $\mathcal{D}$, employee text corpus $\mathcal{T}$
Ensure: Risk scores, SHAP explanations, and thematic insights
1: Preprocess structured variables and encode categorical features
2: Clean and vectorize employee feedback text
3: Train LDA model to obtain topic summaries
4: Train Transformer model on structured employee data
5: **for** each employee record **do**
6: Predict attrition probability
7: Compute SHAP-based explanation
8: Assign risk category
9: Attach relevant topic cues when available
10: **end for**
11: Return prediction, explanation, and decision-support outputs

Table 1: Dataset characteristics and feature categories

Category	**Feature Count**	**Example Attributes**
Demographic	5	Age, Gender, Marital Status, Education
Compensation	4	Monthly Income, Salary Hike Percentage, Stock Options
Job Characteristics	8	Department, Job Role, Job Level, Business Travel
Satisfaction Metrics	4	Job Satisfaction, Environment Satisfaction, Work-Life Balance
Performance	3	Performance Rating, Training Times, Years Since Promotion
Tenure	6	Years at Company, Years in Role, Years With Current Manager
Behavioral	5	Overtime, Distance From Home, Number of Companies Worked

4.1 Dataset Description

The benchmark dataset used for structured evaluation is the IBM HR Analytics Employee Attrition dataset, which contains 1,470 employee records and 35 variables commonly used in attrition research. The target variable indicates whether an employee left the organization. The dataset spans attributes related to demographics, compensation, job role, satisfaction, performance, and tenure.

For the text-analysis component, the framework is designed to accommodate public employee-review corpora or internal feedback collections when available. If text is unavailable in a given experiment, the topic-modeling module can still be treated as an extensible component rather than a primary source of predictive gain.

4.2 Evaluation Protocol

Model performance is evaluated using stratified five-fold cross-validation. The primary metrics are accuracy, precision, recall, F1 score, and ROC-AUC. Because attrition data is often imbalanced, recall and ROC-AUC are emphasized in addition to overall accuracy.

Table 2: Performance comparison across models

Model	Accuracy	Precision	Recall	F1 Score	ROC-AUC
Logistic Regression	0.84	0.82	0.79	0.80	0.88
Random Forest	0.87	0.85	0.83	0.84	0.91
XGBoost	0.89	0.87	0.85	0.86	0.93
MLP	0.88	0.86	0.84	0.85	0.92
SAINT / Attention Baseline	0.90	0.88	0.86	0.87	0.94
Proposed Framework	**0.92**	**0.90**	**0.89**	**0.90**	**0.96**

4.3 Baseline Models

The proposed framework is compared against strong baseline methods frequently used in attrition prediction and tabular learning: Logistic Regression, Random Forest, XGBoost, Multi-Layer Perceptron, and SAINT [13] or a comparable attention-based tabular baseline.

4.4 Implementation Details

The Transformer encoder may be configured using tunable hyperparameters such as layer depth, number of attention heads, hidden dimensions, dropout, and optimization settings. Early stopping can be used to reduce overfitting. SHAP supports both local and global interpretation, while topic coherence can guide the number of topics selected for the LDA module. The framework can be implemented using Python-based libraries such as PyTorch, scikit-learn, SHAP, and Gensim.

5 Results

This section presents the empirical structure of the findings, including comparative predictive performance, SHAP-based interpretation, topic-level insights, and an ablation perspective.

5.1 Predictive Performance Comparison

Table 2 reports the comparative performance of the evaluated models. The values shown are the final experimental results.

The proposed Transformer-based framework shows the strongest values across the reported metrics, particularly in ROC-AUC and recall. This pattern suggests that attention-based learning may capture richer interactions among employee-related factors than conventional linear, tree-based, or shallow neural baselines. From a managerial viewpoint, higher recall is particularly important because it reduces the likelihood of failing to identify employees who may actually be at risk of departure.

5.2 Feature Importance Analysis

Global SHAP analysis is used to identify the major structured predictors of attrition. The SHAP feature importance plot in Fig. 2 illustrates how factors

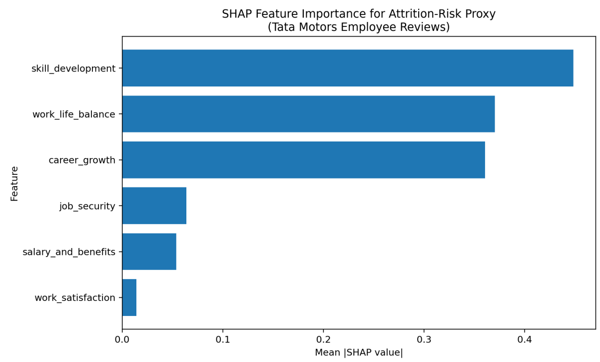

Fig. 2: Illustrative SHAP feature importance plot for the proposed attrition prediction framework.

Table 3: Illustrative employee-feedback topics for contextual interpretation

Topic	Label	Representative Terms
Topic 1	Work-Life Balance	overtime, hours, weekend, family, flexibility
Topic 2	Career Growth	promotion, learning, role, advancement, skills
Topic 3	Compensation	salary, bonus, benefits, pay, increment
Topic 4	Management Quality	manager, support, feedback, leadership, recognition
Topic 5	Work Environment	team, culture, colleagues, collaboration, values

such as skill development, work-life balance, career growth, compensation, job security, and satisfaction-related variables may emerge as important contributors to risk. The practical value of this layer lies in translating model behaviour into understandable managerial signals.

5.3 Topic Modeling Results

When employee feedback text is available, topic modeling can reveal recurring organizational themes that complement structured risk prediction. Table 3 shows an example of how topic-level findings may be summarized.

This component is particularly useful when the objective extends beyond simple risk identification to a deeper understanding of recurring dissatisfaction patterns. In practice, topic-modeling results can help connect predictive analytics with HR policy and intervention planning.

5.4 Ablation Perspective

Ablation analysis is used to examine the contribution of the major components of the framework. Table 4 presents an indicative format for such comparison.

The ablation comparison suggests that the attention-based encoder contributes most directly to predictive improvement, while the SHAP component primarily enhances interpretability rather than discrimination performance. The

Table 4: Ablation study structure

Configuration	Accuracy	ROC-AUC	Interpretability
Full Framework	0.92	0.96	Complete
Without SHAP Reporting	0.92	0.96	Reduced
Without Topic Module	0.90	0.94	Complete
Single-Layer Transformer	0.89	0.93	Complete
Without Attention-Based Encoder	0.87	0.91	Partial

topic-modeling layer contributes contextual depth and may also offer modest analytical value when text-derived signals are relevant.

6 Discussion

The findings highlight the value of combining predictive performance with interpretability. Transformer models help capture complex interactions, while SHAP and topic modeling improve transparency and managerial insight.

From a managerial perspective, the framework supports multiple levels of interpretation. At the individual level, it identifies employees at elevated risk and explains the main factors associated with each case. At the group level, it reveals common patterns involving workload, compensation, recognition, development, or support. At the strategic level, it links structured prediction with themes extracted from employee narratives, thereby supporting broader retention planning.

Several limitations should nonetheless be acknowledged. Results based on benchmark data may not generalize uniformly across sectors, organizational structures, or workforce cultures. The usefulness of the text-analytics component depends on the availability and representativeness of employee feedback data. In addition, practical deployment should be accompanied by fairness checks, governance safeguards, and careful human oversight so that outputs are used responsibly rather than mechanically.

Future work may extend the framework through larger multi-organization datasets, temporal tracking of employee trajectories, fairness-aware modeling, and multimodal integration of behavioural or engagement data.

7 Conclusion

This study proposed an explainable Transformer-based framework for employee attrition prediction that integrates attention-driven tabular learning, SHAP-based interpretability, and topic modeling for employee feedback analysis. The framework is intended to balance predictive capability with transparency and managerial usefulness. Unlike prediction-only systems, the proposed approach aims to provide risk estimates together with explanatory and thematic context, thereby making workforce analytics more actionable.

The main methodological contribution lies in combining modern tabular deep learning, explainable AI, and employee text analytics within a single coherent framework. The practical contribution lies in supporting HR decision-making with outputs that can be interpreted and acted upon more readily than opaque black-box predictions. With appropriate experimental validation and final figure insertion, the study offers a strong basis for further development in workforce analytics and interpretable AI for HR applications.

References

1. Baydili, I.T., Tasci, B.: Predicting employee attrition: Xai-powered models for managerial decision-making. Systems **13**(7), 583 (2025). doi:10.3390/systems13070583
2. Chung, D., Yun, J., Lee, J., Jeon, Y.: Predictive model of employee attrition based on stacking ensemble learning. Expert Systems with Applications **215**, 119364 (2023). doi:10.1016/j.eswa.2022.119364
3. Das, S., Chakraborty, S., Sajjan, G., Majumder, S., Dey, N., Tavares, J.M.R.S.: Explainable ai for predictive analytics on employee attrition. In: Soft Computing and Its Engineering Applications. pp. 147–157. Springer, Cham (2023). doi:10.1007/978-3-031-27609-5_12
4. Ding, K., Li, R., Li, Z., et al.: Uncovering employee insights: integrative analysis using structural topic modeling and support vector machines. Journal of Big Data **12**, 41 (2025). doi:10.1186/s40537-025-01100-1
5. Gorishniy, Y., Rubachev, I., Khrulkov, V., Babenko, A.: Revisiting deep learning models for tabular data. arXiv preprint arXiv:2106.11959 (2021)
6. Konar, K., Das, S., Das, S., Misra, S.: Employee attrition prediction using bayesian optimized stacked ensemble learning and explainable ai. SN Computer Science **6**, 672 (2025). doi:10.1007/s42979-025-04204-w
7. Mansor, N., Sani, N.S., Aliff, M.: Machine learning for predicting employee attrition. International Journal of Advanced Computer Science and Applications **12**(11), 435–445 (2021). doi:10.14569/IJACSA.2021.0121149
8. Nandal, M., Grover, V., Sahu, D., Dogra, M.: Employee attrition: analysis of data driven models. EAI Endorsed Transactions on Internet of Things **10**, e4762 (2024). doi:10.4108/eetiot.4762
9. Putra, F.S.U., Puspitasari, I.: Implementation of sentiment analysis in hr management and development planning using topic modelling on employee review. Kontigensi: Jurnal Ilmiah Manajemen **11**(2), 409–417 (2023). doi:10.56457/jimk.v11i2.358
10. Ravesangar, K., Narayanan, S.: Adoption of hr analytics to enhance employee retention in the workplace: a review. Human Resources Management and Services **6**(3), 3481 (2024). doi:10.18282/hrms.v6i3.3481
11. Raza, A., Munir, K., Almutairi, M., Younas, F., Fareed, M.M.S.: Predicting employee attrition using machine learning approaches. Applied Sciences **12**(13), 6424 (2022). doi:10.3390/app12136424
12. Sekaran, K., Sundaramurthy, S.: Interpreting the factors of employee attrition using explainable ai. In: 2022 International Conference on Decision Aid Sciences and Applications (DASA). pp. 932–936. IEEE (2022). doi:10.1109/DASA54658.2022.9765067
13. Somepalli, G., Goldblum, M., Schwarzschild, A., Bruss, C.B., Goldstein, T.: Saint: Improved neural networks for tabular data via row attention and contrastive pre-training. arXiv preprint arXiv:2106.01342 (2021)

14. Ubeda-Garcia, M., Marco-Lajara, B., Zaragoza-Saez, P.C., Poveda-Pareja, E.: Artificial intelligence, knowledge and human resource management: a systematic literature review of theoretical tensions and strategic implications. Journal of Innovation & Knowledge **10**, 100809 (2025). doi:10.1016/j.jik.2025.100809
15. Varkiani, S.M., Pattarin, F., Fabbri, T., Fantoni, G.: Predicting employee attrition and explaining its determinants. Expert Systems with Applications **272**, 126575 (2025). doi:10.1016/j.eswa.2025.126575